

Problem D

Spam or Not Spam

Time Limit: 5 seconds

Unsolicited email (spam) is annoying and clutters your mailbox. You are to write a spam filter - a program that reads email messages of regular ASCII characters (32-127) and tries to determine whether or not each message is spam.

How can we determine whether or not a message is spam? Spam contains words and phrases that are not common in genuine email messages. For example, the phrase

MAKE MONEY FAST, HONEY!!

is in all-uppercase, contains the word "money" and ends with a double exclamation mark.

One way to create a spam filter is to read through many spam and non-spam messages and to come up with a set of rules that will classify any particular message as spam or not. This process can be tedious and error prone to do manually. Instead you will write a program to automate the process.

A useful step in automatic classification is to split the text up into set of trigrams. A trigram is a sequence of three adjacent characters that appear in the message. A trigram is case sensitive. The example above is composed of the trigrams:

MAK
AKE
KE
E M
MO
MON
ONE
NEY
EY
Y F
FA
FAS
AST
ST,
T,
, H
HO
HON
ONE
NEY
EY!
Y!!

If we examine a sample of spam and non-spam messages we find that some trigrams are more common in spam; whereas others are more common in non-spam. This observation leads to a classification method:

- Examine a sample consisting of a large number of spam messages. Count the number of times that each trigram occurs. In the example above, there are 20 distinct trigrams; the

trigrams ONE and NEY occur twice each and the remaining 18 trigrams occur once each. (Trigrams that do not occur are considered to occur 0 times.) More formally, for each trigram t we compute the frequency $f_{\text{spam}}(t)$ with which it occurs in the sample of spam.

- Examine a sample consisting of a large number of non-spam messages. Compute $f_{\text{non-spam}}(t)$, the frequency with which each trigram t appears in the sample of non-spam.
- For each message to be filtered, compute $f_{\text{message}}(t)$ for each trigram t .
- If f_{message} resembles f_{spam} more closely than it resembles $f_{\text{non-spam}}$ it is determined to be spam; otherwise it is determined to be non-spam.
- A similarity measure determines how closely f_1 and f_2 resemble one another. One of the simplest measures is the cosine measure:

$$\text{similarity}(f_1, f_2) = \frac{\sum_t f_1(t) \cdot f_2(t)}{\sqrt{\sum_t [f_1(t)]^2} \cdot \sqrt{\sum_t [f_2(t)]^2}}$$

Then we say that a message is spam if

$$\text{similarity}(f_{\text{message}}, f_{\text{spam}}) > \text{similarity}(f_{\text{message}}, f_{\text{non-spam}})$$

Input

The input file contains several sets (less than 10) of input. The description of each set is given below.

The first line of each set contains three integers: $s(0 < s < 5)$ the number of sample spam messages to follow; $n(0 < n < 5)$ the number of sample non-spam messages to follow; $c(0 < c < 10)$ the number of messages to be classified as spam or non-spam, based on trigram the trigram frequencies of the sample messages. Each message consists of several lines of text and is terminated by a line containing "**ENDMESSAGE**". This line will not appear elsewhere in the input, and is not considered part of the message. No line has more than 1000 characters.

Input is terminated by a set where $s=0$, $n=0$ and $c=0$. This set should not be processed.

Output

For each set of input your program should output a single line to identify the serial of the input set. The output specification for each set is given below:

For each of the c messages, your program will output two lines. On the first line, output $\text{similarity}(f_{\text{message}}, f_{\text{spam}})$ and $\text{similarity}(f_{\text{message}}, f_{\text{non-spam}})$. On the second line print the classification of the message ("spam" or "non-spam"). Round the numbers to five decimal digits.

For detailed description look at the output for sample input.

When forming trigrams, we never include a new line character. We don't include trigrams that span multiple lines, either.

Sample Input

Output for Sample Input

2 1 1

AAAA

BBBB CCCC

ENDMESSAGE

BBBB

ENDMESSAGE

AAAABBBB

ENDMESSAGE

AAABB

ENDMESSAGE

2 1 2

AAAA

BBBB CCCC

ENDMESSAGE

BBBB

ENDMESSAGE

AAAABBBB

ENDMESSAGE

AAABB

ENDMESSAGE

AAABB

ENDMESSAGE

0 0 0

Set 1:

0.21822 0.73030

non-spam

Set 2:

0.21822 0.73030

non-spam

0.21822 0.73030

non-spam

(Problem Setter: Gordon V. Cormack, University of Waterloo, Canada)